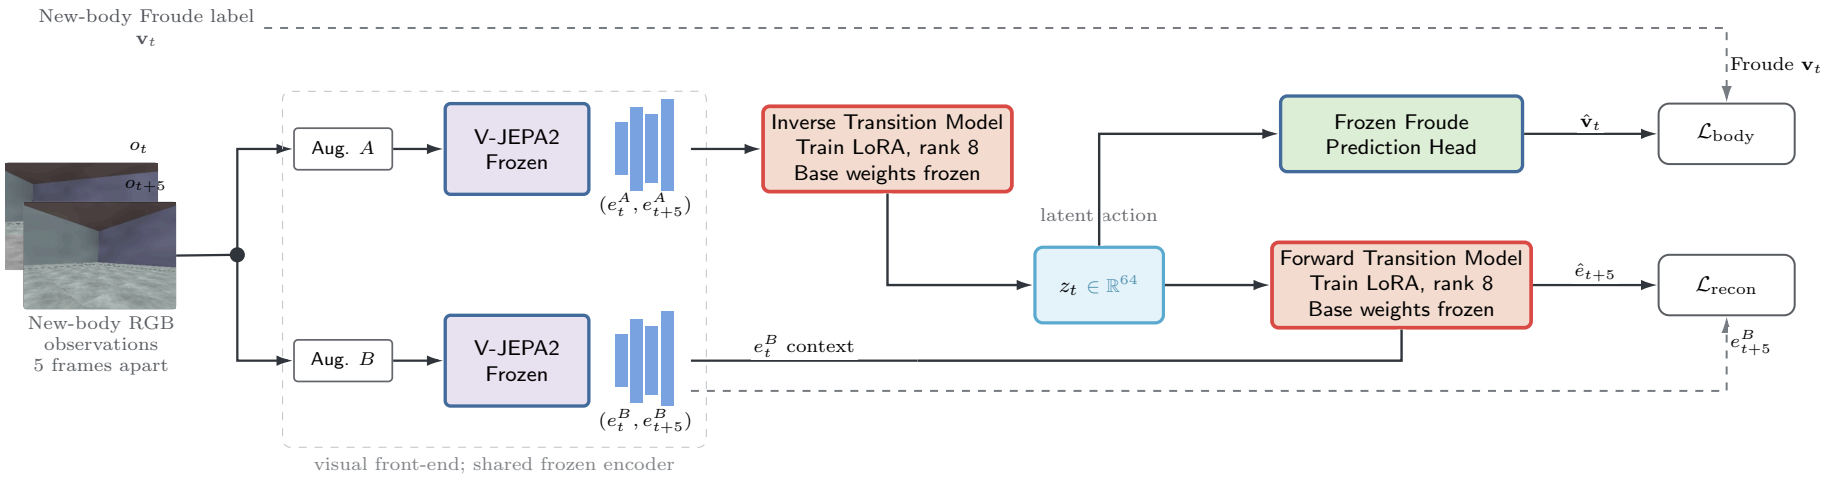
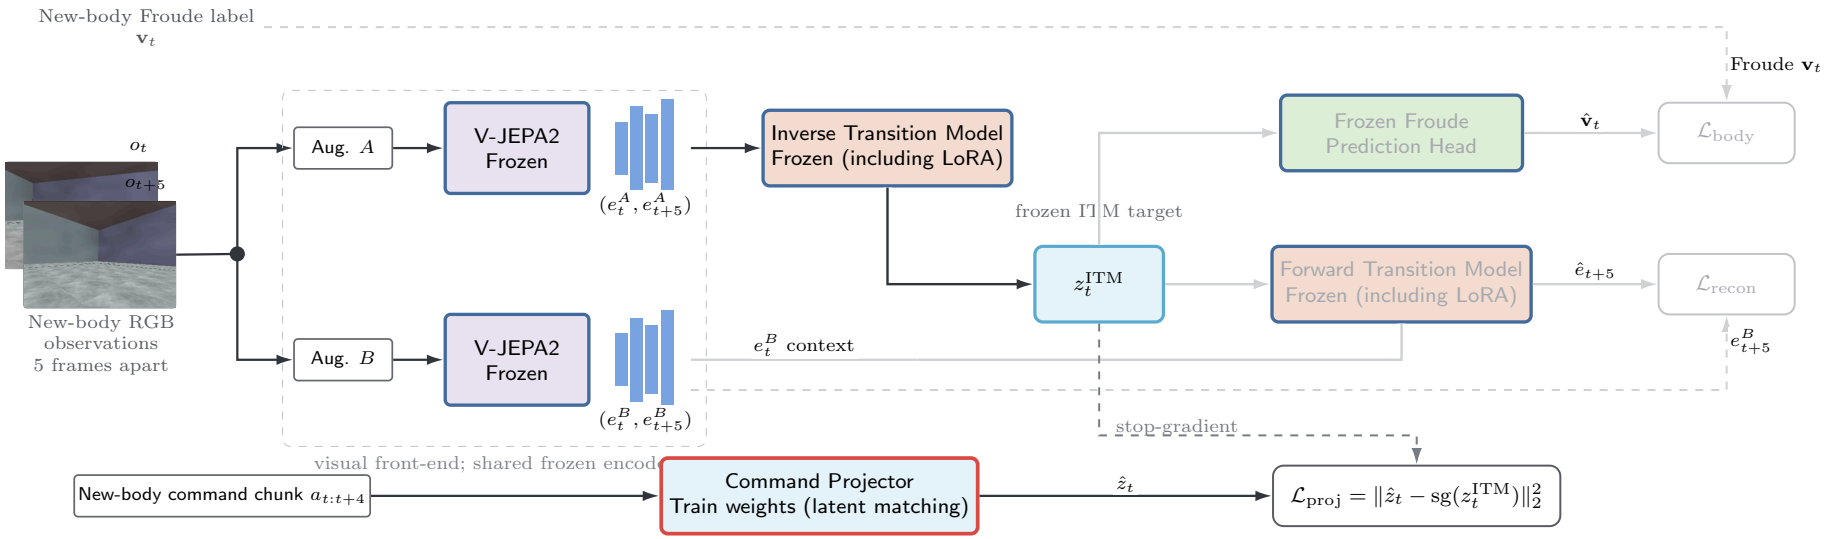
Stage 1: adapt the world model



New body: 44 clips (2.4 min). LoRA adaptation with pretraining losses; gradients pass through the frozen head.

red: trainable; blue border: frozen; dashed: supervision target

Stage 2: fit the new-body command projector



Same new-body clips: supervised latent matching; only the projector receives parameter updates.

red: trainable; blue border: frozen; pale paths: inactive in Stage 2