



Cross-Morphology Locomotion via Visual Latent Action World Models

Disthorn Suttawet

6-month Work Integrated Learning internship
at Vidyasirimedhi Institute of Science and Technology (VISTEC)

Advisor: Mr. Bawornsak Sakulkueakulsuk (FIBO, KMUTT)

Background

Locomotion Controllers Are Developed Around a Body

- Legged locomotion controllers are normally expressed using robot-specific variables: *joint positions, torques, motor velocities, and contact measurements*. These variables work well when the robot body is known and remains unchanged.
- However, **when morphology changes, the same numerical command may produce a different physical outcome.**

Robot bodies do not always remain fixed. Leg dimensions can change between prototypes. Payloads and repairs can alter the dynamics. A controller may also need to operate across multiple related platforms.

A locomotion controller maps the robot's current state to a motor command: $\pi : s_t \rightarrow a_t$

For the simulated stick insect: $a_t \in \mathbb{R}^{18}$

- The learned locomotion policies commonly use proprioceptive states and body-specific actuation.

The study is called **“cross-morphology locomotion transfer”** or **“cross-embodiment generalization”**

$$s_t \xrightarrow{\pi} a_t \xrightarrow[\text{morphology } m]{\text{physical body}} s_{t+1}$$

- So even when morphology is not explicitly given as an input, it is already embedded in the controller through the training process.
- The controller has learned not only “how to walk,” but “how this body walks.”

Background

Controlled Morphology Test

Does morphology changes alone change the locomotion outcome?

Easiest way to prove morphology gap: Same joint command, different body

A joint command is written in the language of motors. But locomotion behavior happens in the language of contact, support, rhythm, balance, and body motion.

- 3 *Medauroidea extradentata* variants in CoppeliaSim, identical in topology and in their 18-dimensional joint action space, differing only in leg length.
- All three received a identical command sequence.

Property	Short	Medium	Long
Leg-length scale	0.5 x	0.75 x	1.0 x
Number of legs	6	6	6
Controlled joints per leg	3	3	3
Action dimension	18	18	18
Command sequence	Identical	Identical	Identical



If the action space is already shared, should the controller transfer automatically?

The null hypothesis is:

H_0 : The locomotion outcomes are equivalent across morphologies.

If this hypothesis cannot be rejected, cross-morphology transfer is not yet a meaningful problem in this experimental setting.

Morphology	Path length	Net displacement
Long, 1.0x	5.217 ± 0.069 (cv 1.3%)	4.404 ± 0.187 (cv 4.2%)
Medium, 0.75x	4.149 ± 0.019 (cv 0.4%)	3.569 ± 0.010 (cv 0.3%)
Short, 0.5x	3.228 ± 0.011 (cv 0.3%)	2.729 ± 0.011 (cv 0.4%)



Background

Need a Middle Language of Locomotion Change

Motor action


$$a_t = \left[q_{1,t}^{\text{target}}, \dots, q_{18,t}^{\text{target}} \right]^{\top}$$

Different bodies may require different motor commands to realize that same behavioral event.


physical locomotion actions

- lifting a foot into swing,
 - placing a foot into stance,
 - pushing the body forward,
 - transferring support between legs,
 - correcting lateral instability.

Same motor command ≠ same physical consequence



long: (-12.3°, -124.3°)
short: (-53.6°, -41.8°)



angles identical → foot positions differ

We need something more transferable than joint commands, but more precise than a task label.

Behavioral grounding	it should describe an observable locomotion change.
Predictive sufficiency	together with the current state, it should help predict what happens next.
Morphology robustness	similar behaviors should remain comparable across different bodies.
Executability	it should retain enough information to recover body-specific motor commands.

- In reinforcement learning, action usually means the numerical output of a policy. *In locomotion, however, what matters physically is not the number itself.* What matters is the resulting transition: a foot is lifted, contact is established, the body moves forward, or balance is recovered.
- Therefore, cross-morphology transfer may not require identical motor commands. It may instead *require preserving the meaning of the locomotion event while allowing its motor realization to change with the body.*

Literature Review

Existing methods transfer by supplying information about the new body

Strategy		Main idea	Body-specific information required	Limitation
Per-body retraining	DreamerV3, Hafner et al. (2023).	Learn a new controller for every morphology	New interactions, reward evaluations, and optimization	Training cost is repeated for each body
Morphology conditioning	QWM, Danesh et al. (2026).	Condition the policy or world model on physical parameters	Leg lengths, masses, torque limits, CAD, URDF, or USD descriptions	Assumes the morphology is accurately known
Online system identification		Infer hidden dynamics from recent interaction history	Proprioception, previous commands, and observed responses	Requires access to the robot's internal signals and sufficient interaction
Shared policy with robot-specific adapters	L3P, Zheng et al. (2025).	Share a common policy backbone while learning separate encoders or decoders	Sensor definitions, joint ordering, actuator conventions, and adaptation data	A new adapter must still be fitted for each platform
Common structure		Although their implementations differ, these approaches obtain information about the new body through at least one of three channels Design specification ∨ Internal sensing ∨ New interaction		

What happens when the body can be observed moving, but cannot be accurately described in advance?

This might occur with animals, undocumented hardware, repaired robots, payload changes, structural wear, or systems whose real physical parameters no longer match their design files.

These methods are effective when the new body can be accurately described or instrumented.
Their common assumption is: The learner has access to an internal description of the body.

Literature Review

Research Gap: The Body Is Observable but Not Described

proprioception does not provide a naturally shared observation space across substantially different bodies.

Existing approaches assume that the new embodiment can be accessed internally

The previous strategies work when at least one of the following is available:

- an accurate geometric and dynamic model,
- known joint and actuator conventions,
- proprioceptive or force measurements,
- permission to collect new interaction data,
- sufficient time to retrain or adapt.

Research gap

Can visual transitions provide a shared behavior-level representation across morphology changes, and how does that representation compare with proprioception?

- For a laboratory robot. CAD, URDF, motor specifications, and onboard sensors is still an approximation.
- The documented body is the set of parameters we believe the system has. The physical body is the system that actually moves and interacts with the environment. In practice, these two descriptions may differ.
- Motivating use cases: Animals or Undocumented machines.
- We may observe how the system walks, while having no access to its joint encoders, actuator commands, or exact dynamic parameters.
- Across substantially different embodiments, proprioceptive vectors may differ in dimension, ordering, units, and semantics. External vision offers a common observation format.

If the body cannot be described directly, perhaps its dynamics can be inferred from how its observable state changes over time.

Literature Review

Why Test Vision When Proprioception Is Available?

Most locomotion systems use proprioception:

- joint angles
 - joint velocities
 - IMU
 - foot forces
 - motor states
- These are powerful, but they are body-internal and convention-specific.

Vision is different:

- same input format for different bodies
- external observation
- does not require joint conventions or CAD
- captures what happened in the world

But vision is also dangerous:

A camera sees morphology very clearly.
In the current pilot, raw V-JEPA2 features
decode morphology at around 99%.

Can the latent action model extract behavior-relevant change from visual features that strongly contain body shape?

This thesis does not assume vision is superior. It tests whether vision can recover a transferable behavior representation and compares it against proprioception.



DIFFERENT INTERNAL INTERFACES

different dimension, and a different semantic ordering of the same quantities



COMMON EXTERNAL FORMAT

same tensor shape on any body, obtainable without access to the body's interior

Common format is a precondition, not the result: morphology still decodes from raw features at 99.9%.
Removing that is what the inverse transition model is for.

Problem Formulation

External observation shows the consequences of hidden body dynamics

Suppose an observer records a sequence:

$$O_1, O_2, O_3, \dots, O_T$$

where O_t is an external observation of the body at time t .

A single observation may reveal:

- body shape,
- limb configuration,
- approximate pose,
- foot locations,
- surrounding terrain.

External observation does not give us the robot's exact link lengths, masses, actuator limits, or internal sensor readings. But it gives us something else: *the visible consequences of those hidden properties*.

A single image mainly tells us what the system looks like. A transition between images tells us how the system behaves.

These changes reveal:

- which limbs are moving,
- which feet enter or leave contact,
- how the body responds to support,
- whether the body progresses or slips,
- how balance changes over time.

$$O_t, a_t \rightarrow O_{t+1}$$

Therefore, the learning problem becomes:

$$s_{t+1} = f_m(s_t, a_t)$$

Given the current observation and some representation of the locomotion action, can we predict the next observation?

The morphology m may be unknown, but it influences the observable transition:

Can the model learn the relevant dynamics from observed state transitions?

This changes the problem from body description to transition prediction.

Literature Review

Why a World Model?

A world model is often described as a learned simulator used for planning or reinforcement learning. But the important feature for this thesis is that a world model must predict consequences.

A world model learns how actions produce state transitions by approximating the dynamics of the environment

$$\hat{s}_{t+1} = F_{\theta}(s_t, a_t)$$

For visual observations, images are first converted into compact representations:

$$e_t = E(o_t)$$

The transition model then predicts:


$$\hat{e}_{t+1} = F_{\theta}(e_t, a_t)$$

The model is trained by comparing the prediction with the actual next observation:

$$\mathcal{L}_{\text{prediction}} = \|\hat{e}_{t+1} - e_{t+1}\|_2^2$$


The world model is not introduced primarily to generate imagined rollouts. It is introduced because transition prediction provides an objective test of whether an action representation captures meaningful locomotion change.

(Recurrent State Space Model tested with gameplays with explicit action label)



Mastering Diverse Domains through World Models (Google DeepMind 2023)

Danijar Hafner, Jurgis Pasukonis, Jimmy Ba, Timothy Lillicrap



(Cross-morphology do not share the same action)

DreamerV3 (Hafner et al., 2023) learns a compact latent state and imagines rollouts inside it, so the policy trains without touching the environment. One algorithm with one hyperparameter set covers more than 150 domains. The limitation for this project: each domain still requires its own world model. Nothing carries across bodies, and the method assumes explicit action labels throughout.

A conventional world model still conditions its prediction on the robot's native motor command $\backslash(a_t\backslash)$. But we have already shown that the physical meaning of $\backslash(a_t\backslash)$ changes with morphology. Therefore, a standard world model reproduces the same body-specific action problem inside its transition model.

Literature Review

What Should the World Model Treat as an Action?

The standard formulation is:

$$(e_t, a_t) \rightarrow e_{t+1}$$

But across morphologies:

$$a_t = \text{same numerical command}$$

Therefore, instead of assuming that the native joint command is the shared action representation, we ask whether an intermediate variable can be inferred from the observed transition:

$$e_{t+1} - e_t = \text{different physical change}$$


We propose $\{z_t\}$ as a candidate middle language: a compact representation inferred from the locomotion change between consecutive observations.

$$(e_t, e_{t+1}) \rightarrow z_t$$



LATENT ACTION ROBOT FOUNDATION WORLD MODELS FOR CROSS-EMBODIMENT ADAPTATION

Anonymous



LAC-WM (Huang et al., ICML 2026) discards explicit action labels as the conditioning signal. An inverse model infers an abstract action z from consecutive observations, and the world model is conditioned on z rather than on any robot's native command vector. One latent space then covers several embodiments, and adding embodiments improves it rather than fragmenting it.

- LAC-WM problem: different robots have different action formats
- This project: robots share the same 18D action format, but the same 18D command has different dynamics meaning



Research Question

Can observable locomotion change become a shared action representation?

Primary research question

Can a latent action inferred from visual state transitions preserve behavior-relevant locomotion information while reducing morphology-specific information across different leg lengths?

Secondary question

How does the visual latent representation compare with raw visual features, native joint commands, and proprioceptive observations?

Testable hypotheses	
H1 : Behavioral information	The latent action should preserve information about locomotion events such as foot contact and support transition:
H2: Reduced morphology dependence	The morphology should be less recoverable from $\mathbf{z}_t$ than from the raw visual representation:
H3: Predictive isefulness	The current state and latent action should predict the next state:
H4: Added value over native commands	A latent-conditioned dynamics model should be compared against a model conditioned directly on the 18D joint command. (If the native-command model performs equally well across morphologies, the latent action may be unnecessary in this controlled setting.)

***Success is not defined as “vision wins”**
The intended result is a scientific comparison: Determine which information source produces the most transferable behavior representation, under which assumptions, and at what cost.

●

Experimental Design

Can observable locomotion change become a shared action representation?

Execution Plan

Platform:

- CoppeliaSim

Robot:

- Medauroidea extradentata stick-insect model
- 3 morphologies (short / medium / long)
- each with 6 legs, 3 controlled joints/leg

Native action:

- joint position targets $\mathbb{R}^{18}$

Data collection policy: TBD

Behaviors:

- walk (might extent to turn / stop)

Train morphologies: Medium leg is an interpolation test set (Extrapolation future direction)

- short + long leg


Test morphology:

- medium leg (Step 2 transfer test)

Camera Control

- Script-created camera identical across morphologies.
- Third-person view.
- Approximately 40° elevation.
- 45° field of view.
- Matte, lightly textured floor.
- Camera follows the body from a fixed relative offset.
- No empty background pixels in the validated configuration.

Validation Logic





Evaluation

What Counts as Success?

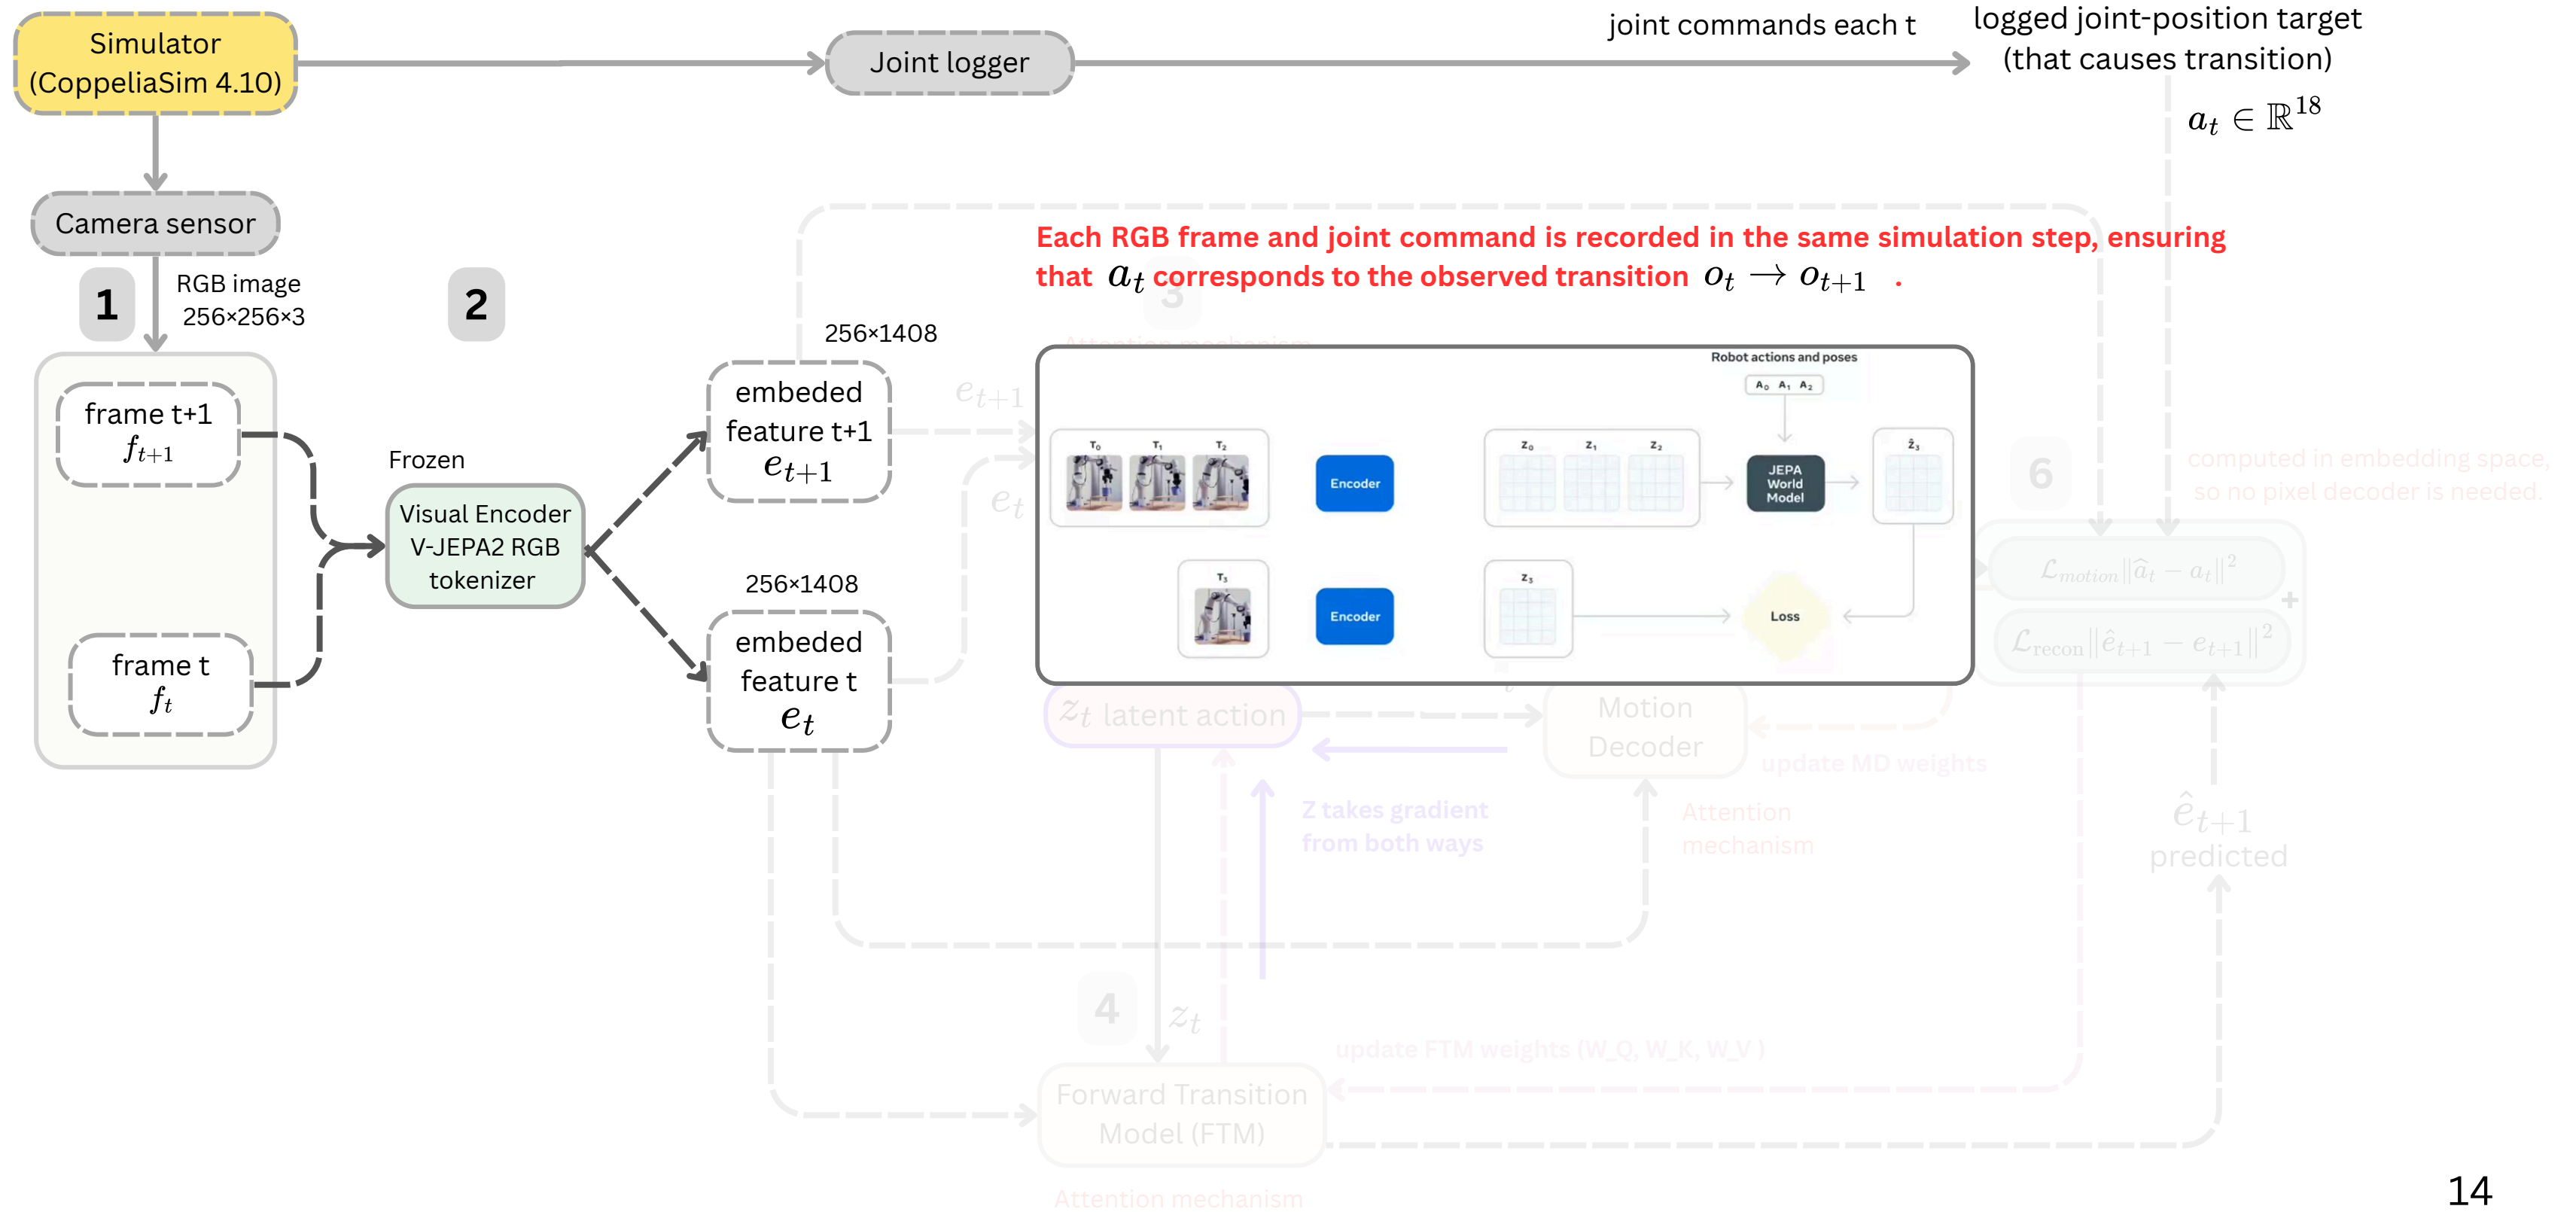
Requirement	Measurement	Success criterion
Behavior preservation	Foot-contact balanced accuracy and macro-F1	$\mathcal{B}(z_t)$ transfers better than raw visual features $\mathcal{B}(e_t)$
Morphology robustness	Morphology probe accuracy; silhouette as supporting evidence	Morphology is less decodable from $\mathcal{B}(z_t)$ than from $\mathcal{B}(e_t)$
Predictive value	Next-embedding error on held-out medium	$\mathcal{F}(e_t, z_t)$ transfers better than or degrades less than $\mathcal{F}(e_t, a_t)$
Executability	IK command reconstruction MAE	$\mathcal{MD}(e_t, z_t)$ outperforms $\mathcal{MD}(e_t, 0)$
Adaptation efficiency	Medium-leg learning curve	Pretrained model reaches the target error using fewer episodes than training from scratch

$$\text{BehaviorTransfer}(z_t) > \text{BehaviorTransfer}(e_t)$$
$$\text{MorphologyDecode}(z_t) < \text{MorphologyDecode}(e_t)$$

Both conditions must hold simultaneously. Reducing morphology information alone is insufficient because a collapsed representation would also make morphology difficult to decode.


Pipeline (Focused)

Data source and Pre-processing



Implemented Front-End V-JEPA2

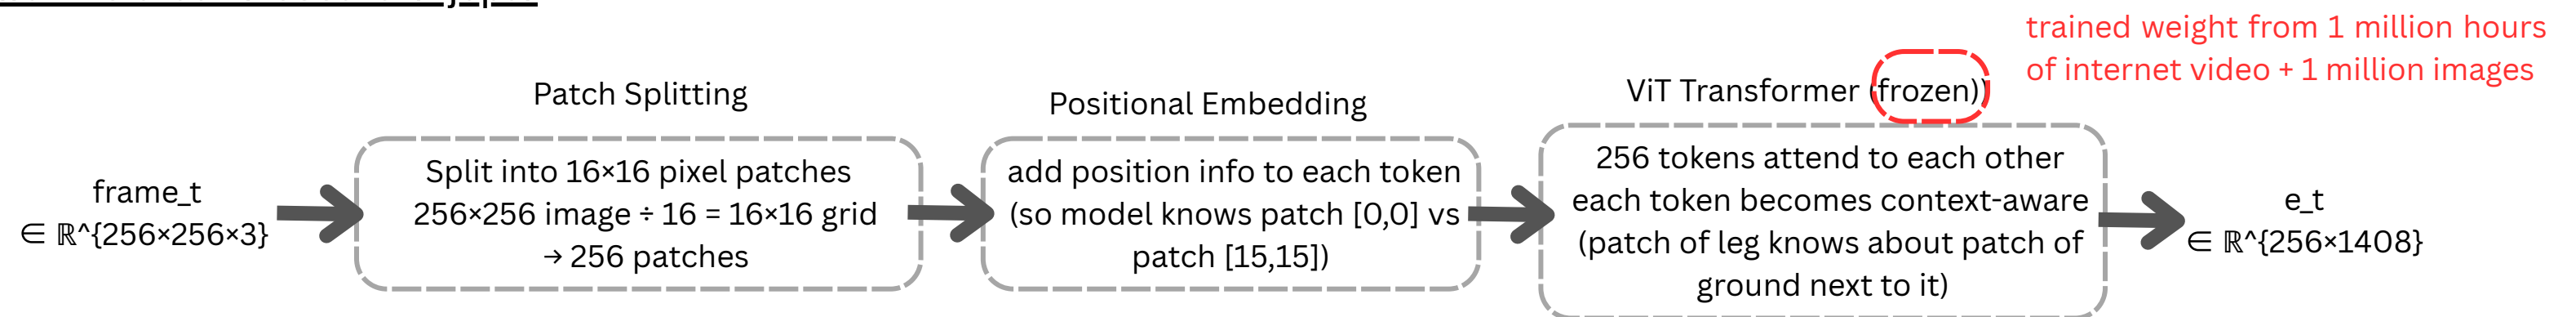
V-JEPA2 is self-supervised on roughly one million hours of video, with an objective that predicts masked content in representation space rather than pixel space. That objective rewards motion-relevant features, which is why it is a reasonable starting point for gait.



<https://ai.meta.com/research/vjepa/>

We adopt V-JEPA 2's frozen RGB tokenizer as our visual encoder.

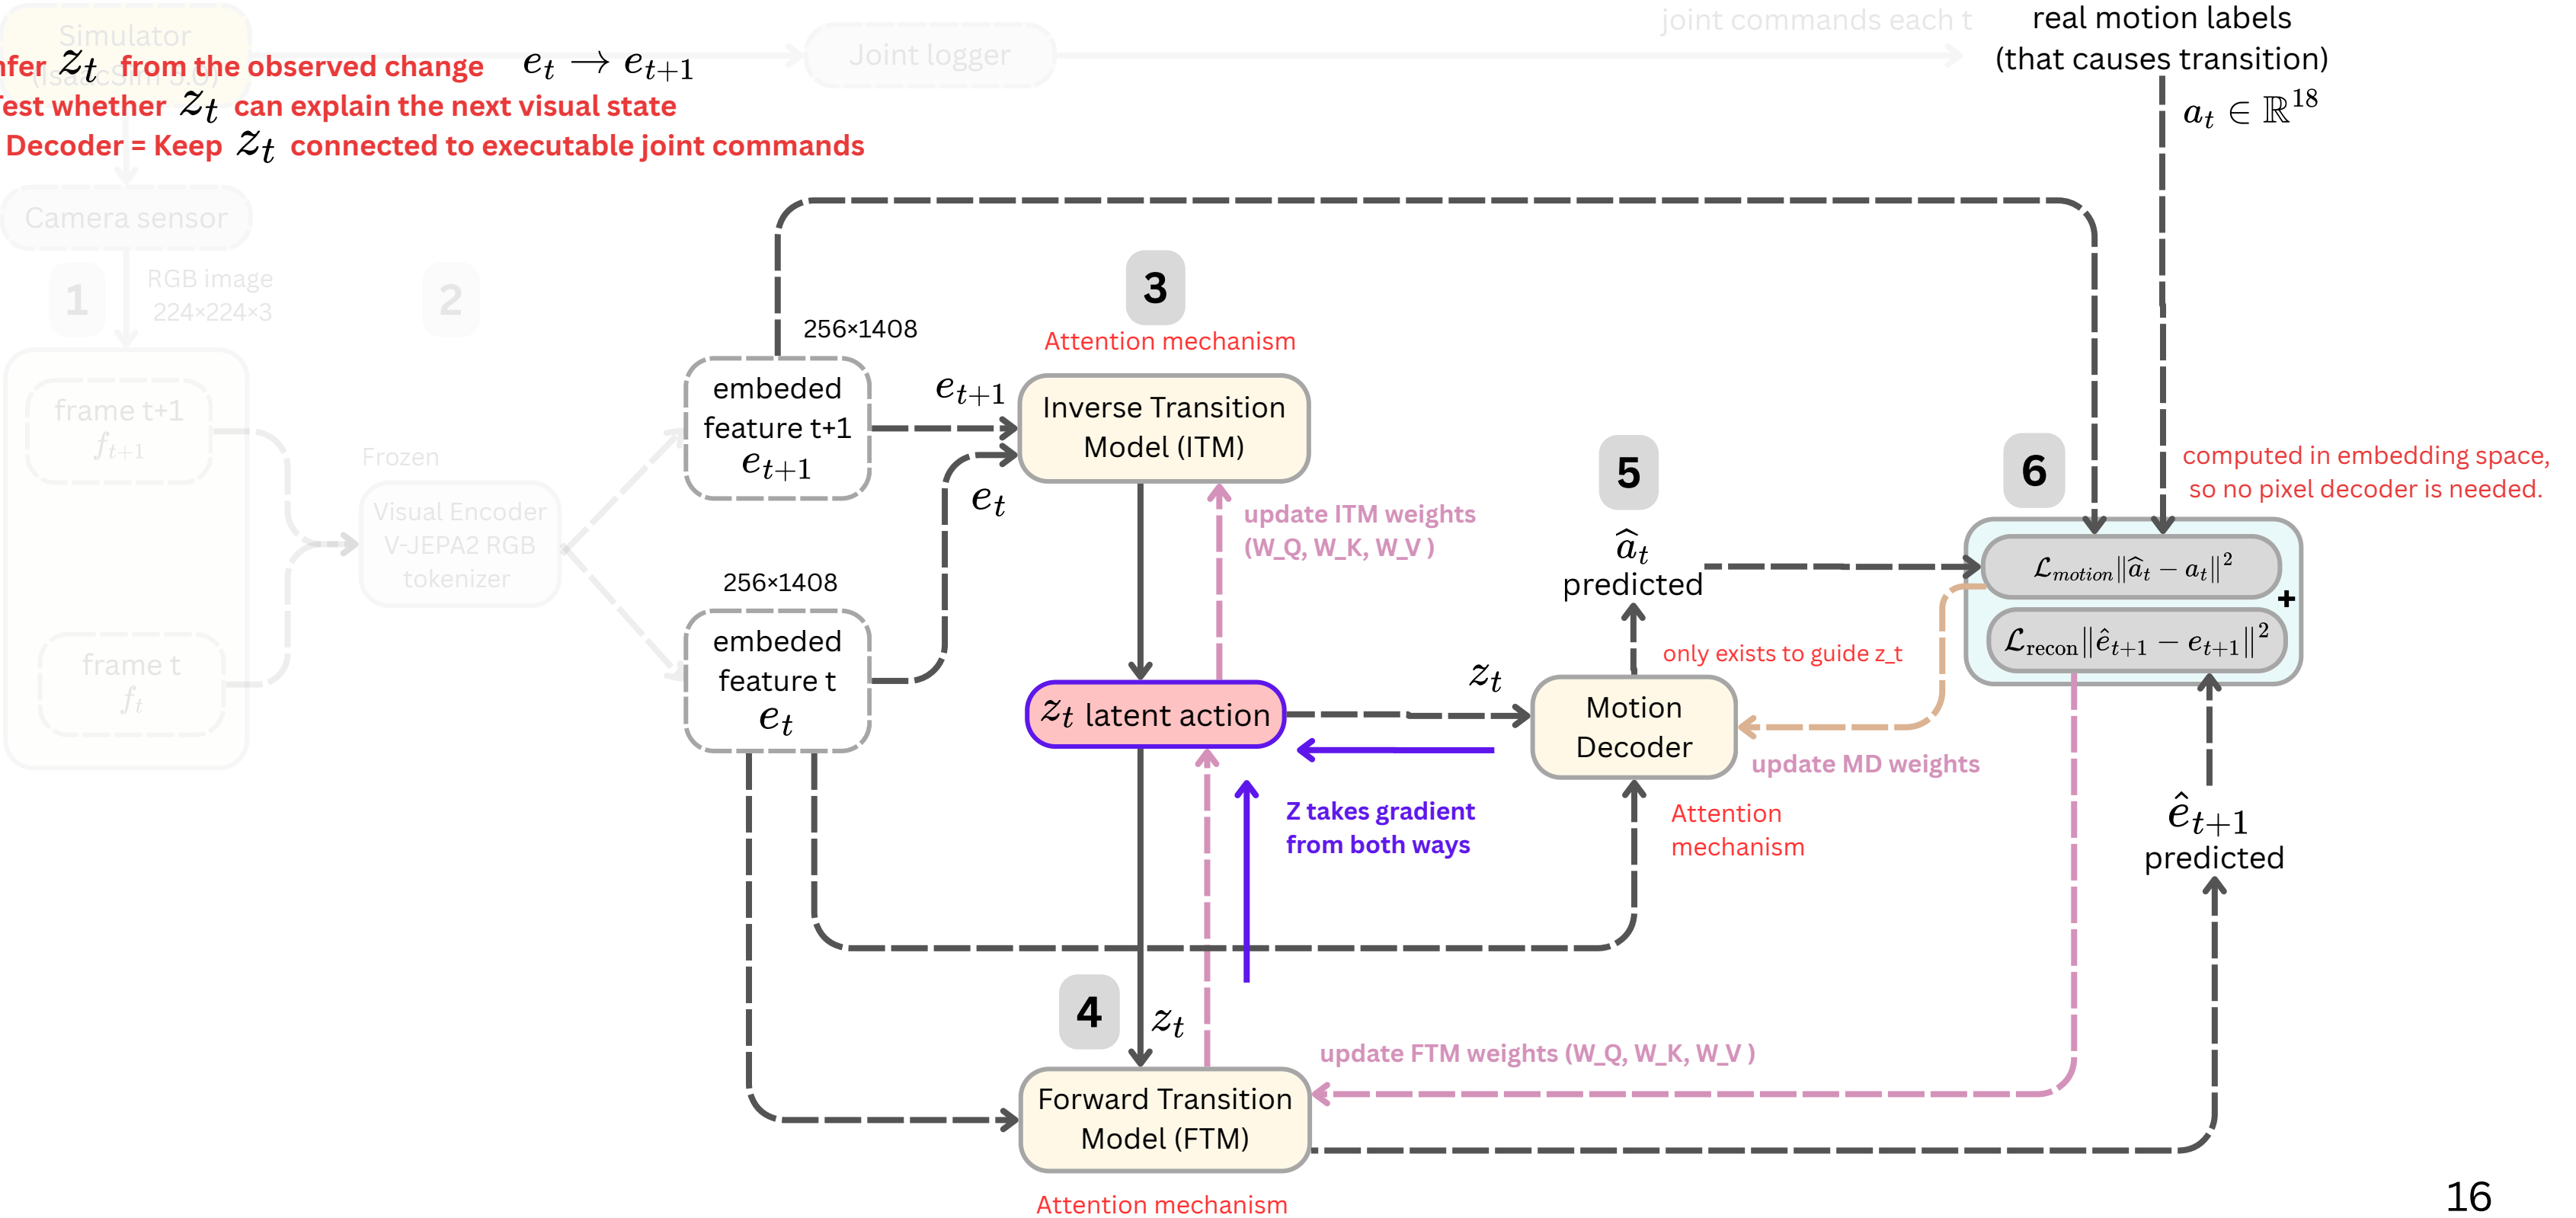
- extracting per-frame embeddings $e_t \in \mathbb{R}^{256 \times 1408}$ from simulation video without any fine-tuning.
- These embeddings serve as input to our ITM and FTM, which learn a morphology-agnostic latent action space z_t shared across stick insect morphologies.
- V-JEPA2 is a reasonable starting point because it was pretrained on large-scale video; whether its features contain locomotion information is tested empirically in Step 0



Pipeline (Focused)

How is z_t learned

ITM = Infer z_t from the observed change $e_t \rightarrow e_{t+1}$
 FTM = Test whether z_t can explain the next visual state
 Motion Decoder = Keep z_t connected to executable joint commands



Pipeline (Focused)

Proposed Trainable Modules

Compresses the observed transition into a latent action that explains what changed between the two states.

ITM (Inverse State-Transition Model) $(e_t, e_{t+1}) \rightarrow z_t$

: $[e_{t^1}, e_{\{t+1\}^1}] \in \mathbb{R}^{\{512 \times 1408\}} + q_t \in \mathbb{R}^{\{512\}}$

- $\times 4$ Causal Self-Attention $\rightarrow e_{\{t+1\}}$ absorbs context of e_t
- $\times N$ Cross-Attention $\rightarrow q_t$ queries S , focuses on moving patches

Output: $z_t \in \mathbb{R}^{\{64\}}$ (latent action)

z_t latent action

FTM (Forward State-Transition Model) $(e_t, z_t) \rightarrow \hat{e}_{t+1}$

: $e_{t^2} \in \mathbb{R}^{\{256 \times 1408\}} + z_t \in \mathbb{R}^{\{64\}}$

- $\times 8$ Transformer Blocks: self-attn(e_{t^2}) $\rightarrow$ self-attn(z_t) $\rightarrow$ cross-attn(e_{t^2} queries z_t)

Output: $\hat{e}_{\{t+1\}} \in \mathbb{R}^{\{256 \times 1408\}}$
Loss: $L_{\text{recon}} = ||\hat{e}_{\{t+1\}} - e_{\{t+1\}^1}||^2$

Tests whether the current state and z_t are sufficient to explain the next visual state.

Motion Decoder $(e_t, z_t) \rightarrow \hat{a}_t$

Input: $z_t \in \mathbb{R}^{\{64\}} + e_{t^1} \in \mathbb{R}^{\{256 \times 1408\}}$

- cross-attn(z_t queries e_{t^1}) + MLP

Output: $\hat{a}_t \in \mathbb{R}^{\{18\}}$
Loss: $L_{\text{motion}} = ||\hat{a}_t - a_t||^2$
(drop after pretraining finish)

Requires z_t to retain information connected to executable, body-specific joint commands.

$L_{\text{total}} = L_{\text{recon}} + \lambda \cdot L_{\text{motion}}$

ITM receives training signals from both objectives through z_t .


Pipeline (Focused)

What prevents a shortcut?

In the previous pipeline, z_t may become a compressed copy of the next state rather than a representation of the action between states.

Fix

Use the same augmentation parameters for f_t and f_{t+1} inside each branch, so the temporal change is preserved. However, Augmentation A and B must be independently sampled.



Preliminary Check

Frozen V-JEPA2 Features Encode Morphology/Session Identity


Experimental Question: Does the frozen visual representation e_t contain locomotion information, and how strongly does it contain morphology/session identity?

Input: frozen V-JEPA2 e_t (1408-d per frame (ViT-g/16, mean-pooled over 256 patches, encoder never trained))

Model: logistic regression (linear probe) + standardization

Data: 3 bodies $\times$ 5 episodes $\times$ 200 steps = 3000 frames, same bit-identical command sequence

Evaluation: 5-fold cross-validation



Morphology Encoding

Target: which body (long / medium / short) — 3 classes, chance 33%

Three tests, weakest $\rightarrow$ strongest:

- (a) supervised probe $\rightarrow$ ~100%, holds under episode-grouped CV (not frame memorisation)
- (b) unsupervised PCA $\rightarrow$ PC1 orders short < medium < long with no labels (self-organises)
- (c) UMAP $\rightarrow$ illustration only, not evidence

Reads as: body identity is a dominant axis of e_t .

This is the baseline z_t must reduce.

Behavior Encoding

Target: which feet are planted (foot-contact pattern, top-8) — macro-F1, chance 0.125

Three splits:


within-body train & test same body $\rightarrow$ 0.84 (contact is decodable)

cross-body train one body, test another $\rightarrow$ 0.16 (does not transfer)

shuffle permuted labels $\rightarrow$ $\approx$ chance (confirms 0.84 is real, not an artefact)

Reads as: e_t encodes behavior, but entangled with body shape (the wider the leg-length gap, the worse it transfers (L $\leftrightarrow$ S worst)).


This is the signal z_t must keep while making it body-independent.



Preliminary Check

Pilot Limitation: Identical Commands, Different Physical States

Identical commands produce different contact states. The difference is absent from the command but present in the resulting visual and proprioceptive state.



Type: Direct measurement, reads the simulator state at one timestep

Controlled: the 18D joint command q_t — bit-identical across all 3 bodies

Measured: per-foot contact force (N) from the sim's force sensors, and the rendered RGB frame


Varied: only leg length (1.0 / 0.75 / 0.5x)

Frame shown: long/medium/short_ep0, step 25

Same command → different physical state:

left-middle foot carries 5.7 N (long) / 9.3 N (short) but only 0.3 N (medium) — airborne on the medium body while planted on the other two.

The command cannot tell the bodies apart — it is identical, yet the outcome is not. This is why the shared a_t makes the latent action vacuous, and why per-body commands (IK retargeting) are needed before training.




The Motion Decoder cannot demonstrate morphology-specific command recovery !

Preliminary Check

Better Expert Data: Shared Targets, Body-Specific Commands

shared task-space target → body-specific joint commands



Type: Concept forward/inverse kinematics
Bodies illustrated: link length 1.0× and 0.5×

Foot position depends on angle AND link length. Fixing angles does not fix behaviour; fixing the foot target does — and the per-body angle difference is exactly the body-specific command the Motion Decoder must learn to produce.

- old CSV = joint angles + contact only
Problem: foot trajectory must be derived via forward kinematics).
- new 66k CSV = foot trajectories logged directly
Provide: ready-made IK targets, plus sim binary contact.

Final Experimental Protocol

Final Experimental Protocol: Train on Two Bodies, Test on One

Execution Plan

Platform:

- CoppeliaSim

Robot:

- Medauroidea extradentata stick-insect model
- 3 morphologies (short / medium / long)
- each with 6 legs, 3 controlled joints/leg

Native action:

- task-space foot trajectories

Data collection policy: IK-retargeted expert dataset

Behaviors:

- walk (might extent to turn / stop)

Train morphologies:

- short + long leg

Test morphology:

- medium leg (Step 2 transfer test)

Camera Control

- Script-created camera identical across morphologies.
- Third-person view.
- Approximately 40° elevation.
- 45° field of view.
- Matte, lightly textured floor.
- Camera follows the body from a fixed relative offset.
- No empty background pixels in the validated configuration.

Validation Logic

- ✓

Step 0: Does a morphology gap exist?
- ✓

Step 1: Does frozen vision contain locomotion information?
- ↓

Step 2: Does z_t preserve behavior while reducing morphology information?

Step 3: Does z_t improve held-out transfer over raw commands?

what model	SHORT 0.5× TRAIN 				
	Medium 0.75× TEST 				
what is recorded	Long 1.00× TRAIN 				
	RGB	Proprioception	Command	Contact	Next RGB
Step 0 → Step 1 → Step 2 → Step 3					



Decisive Ablation

Does Latent Action Add Value Beyond Raw Joint Commands?

Compare models under the same encoder, FTM capacity, dataset, and training budget:

Transition model	Conditioning input	What it tests
Observation only	$\mathcal{F}(e_t, 0)$	How predictable is motion without an action?
Raw action	$\mathcal{F}(e_t, a_t)$	Standard body-specific command representation
Latent action	$\mathcal{F}(e_t, z_t)$	Proposed transition-based representation

Primary comparison $\mathcal{L}_{pred}^{medium}(e_t, z_t) < \mathcal{L}_{pred}^{medium}(e_t, a_t)$

The latent representation is useful only if it improves held-out transition prediction or adaptation relative to raw commands.

Possible Outcomes

What Would the Results Mean?

Behavior in z_t	Morphology in z_t	Prediction/transfer	Interpretation
Preserved	Reduced	Improved	Intended morphology-
Preserved	Still high	Improved	Useful representation,
Lost	Reduced	Poor	Representation collapse or
Preserved	Reduced	No improvement	Clean latent space, but no
Preserved	Reduced	Worse than a_t	Raw commands remain more

Success requires all three:

behavior preserved

+

morphology reduced

+

transfer improved

Scope and Limitations

Included

- Simulation-based hexapod locomotion
- Three leg-length morphologies
- Fixed third-person RGB camera
- Forward locomotion and foot-contact behavior
- Medium-leg interpolation test
- ITM, FTM, Motion Decoder, and frozen V-JEPA2 encoder

Not claimed

- Real-robot or animal transfer
- Extrapolation beyond the training morphology range
- Camera-viewpoint invariance
- Generalization to manipulation or complex terrain
- Fully autonomous control through $\mathbf{z}_t$

IK creates comparable task-space objectives, but it does not guarantee identical contact dynamics or identical behavior across morphologies.

Contributions and Milestones

Expected Contributions and Research Milestones

Expected contributions

- A controlled benchmark for studying locomotion representation across changing body morphology.
- A visual latent-action world model combining ITM, FTM, and command reconstruction.
- An evaluation framework separating behavior preservation, morphology leakage, predictive sufficiency, executability, and adaptation efficiency.
- Evidence establishing whether visual latent actions improve held-out morphology transfer over raw joint commands.

Objectives

- To design and train a Latent Action Space pipeline using Inverse State-Transition Model (ITM), Forward State-Transition Model (FTM), and Motion Decoder on video data and Action Labels automatically logged by simulation, from a Stick Insect robot across 3 morphologies.
- To prove that the learned Latent Action is Morphology-Agnostic by using Principal Component Analysis (PCA) to analyze whether z_t from robots of different morphologies performing the same behavior cluster together.
- To test knowledge transfer to an unseen morphology (Medium Leg) not seen during Pretraining, and ***measure whether the resulting World Model reduces the amount of data required for a new morphology.***

Stage	Completion criterion
Dataset redesign	Valid IK gait, synchronized observations, distinct per-body commands
Representation training	Stable ITM–FTM–MD optimization without collapse
Representation evaluation	Behavior retained while morphology/session signal decreases
Transfer evaluation	Held-out medium results compared with required baselines
Thesis conclusion	Determine where latent action helps, fails, or requires stronger supervision

The contribution is not an assumption that latent action must work, but a controlled test of whether transition-based visual representations provide measurable cross-morphology value.